Stance Detection in Tweets

W266 Natural Language Processing - Final Project

Alex Dessouky & Tim Spittle

December 7, 2019

Abstract

TBD - Abstract goes here

1 Introduction

1.1 Background

Stance classification is a natural language processing (NLP) task that seeks to automatically determine the "stance" of a statement in relation to a "target". The stance the author of that statement can take is either in favor of the target, against the target, or neutral toward the target. A target can be anything from a proposition to an idea, a person, or a political issue, etc. Since the statement may not address the target at all, it is possible that the author's stance towards the target of evaluation cannot be determined.

Stance classification of tweets has applications across many domains: marketing industry efforts to measure public opinions on products, political campaign attempts to measure public views on candidates' policies, and efforts by Twitter to identify bad actors (i.e., "trolls"). While there has been significant research in stance classification with respect to debates and online forums, Twitter poses a new challenge as many "tweeters" represent their stance towards a target implicitly and often use figurative language, shorthand, acronyms, or hashtags. Given an abundance of online discourse on social media, automatic stance detection of social media content is an appealing task with widespread potential applications. In particular, Twitter is a social media platform with an incredibly fast-paced environment where many tweets are shared back-and-forth addressing a variety of topics and issues.

1.2 Objective

The objective of this paper is to explore new methodologies for detecting stance of tweets and attempt to design a classification model that outperforms existing models.

2 Data, Task, & Evaluation

We relied on an existing, pre-labeled Twitter dataset that has been analyzed by competitor models against which we can compare our results (Mohammad et al. (2016a)). This dataset was originally compiled for the International Workshop on Semantic Evaluation (SemEval-2016) where it was used in an evaluation exercise, Task 6 (*Mohammad et al. (2016b)*).

The dataset consists of approximately 5,000 tweet-target pairs annotated for both stance and sentiment. The targets may or may not be referred to in the tweets. While sentiment is shown to be beneficial for stance classification, it was not available to competitors at the SemEval-2016 conference. Given this in combination with the fact that pre-labeled sentiment would not be available in a production environment, the sentiment labels were ignored for our classification model (*Sobhani, Mohammad, and Kiritchenko (2016)*).

The process for collecting and annotating these tweets is detailed in *Mohammad et al. (2016a)* . The authors queried Twitter for topic-specific hashtags in an effort to collect a full corpus of tweets related to the intended targets. These targets include:

- Atheism

- Climate Change is a Real Concern

- Feminist Movement

- Hillary Clinton

- Legalization of Abortion

After extracting the results, the authors retained only tweets with the query hashtag at the end. The query hashtag was then removed in order to exclude obvious clues for the classification task. Note that while this is a necessary step for the NLP task, removal of the query hashtag may actually change the stance of the tweet (e.g. if the hashtag was a negation of the preceding text).

Tweets were then annotated for stance using CrowdFlower (http://www.CrowdFlower.com)[http://www.crowdflower.com], a website for crowd-sourced data annotation, where each tweet was annotated by at least eight reviewers. The raw inter-annotater agreement was 73.1%, however a cutoff was applied where only tweets with greater than 60% agreement were retained, increasing the average agreement to 81.85%. Lastly, the resulting tweets were split into train and test sets chronologically. (See the **Limitations** section for more on how this process may affect the classifcation task).

Models were trained independently for each target. Performance was evaluated for each target using a modified F1 score, which took the macro-average of the "favor" and "against" class F1 scores only. This method was chosen to treat the "none" stance as a class that is not of interest (or negative class), though it still will affect the macro-average F1 scores. *[NEED Information Retrieval Source]* For evaluation of model performance across all targets, the micro-average of F1 scores was evaluated (referred to as **F-microT**). This was the score used as the official competition metric. Alternatively, the macro-average of the F1 scores can also be calculated (referred to as **F-macroT**) by averaging the modified F1 scores acheived for each target (*Mohammad et al. (2016b)*).

3 Related Work

The baseline model for the SemEval-2016 Task 6 competition was a linear-kernel Support Vector Machine (SVM) classifier, with 3 different specifications and features sets used (*(Mohammad et al. 2016b)*):

- **Unigrams**: five SVM classifiers (one per target) trained using unigram features

- **N-grams**: five SVM classifiers (one per target) trained using word n-grams of 1-, 2-, and 3-gram length and character n-grams of 2-, 3-, 4-, and 5-gram length [*best performer*]
- **N-grams-combined**: one SVM classifier trained on the combined data for all targets, using the same n-gram features as described above

SVM was chosen as a baseline since it is effective on text categorization tasks and robust on large feature spaces. *Sobhani, Mohammad, and Kiritchenko (2016)* expanded on the n-gram model to also include word embeddings and sentiment lexicon features. Note that the word embeddings in this model were trained using the "full corpus" of nearly 2 million tweets extracted by the authors that compiled the original dataset using query hashtags. These tweets were not provided with the data, therefore the strength of this model may be artificially inflated through the context these word embeddings were able to capture by virtue of being trained with exclusive tweets that align precisely with the task at hand.

Several of the competing teams within the SemEval-2016 Task 6 competition took a 'neural network' approach to the challenge. *Zarrella and Marsh (2016)*, the winner of the competition, leveraged transfer learning to perform the classification task. The team sampled 200 million + relevant tweets to create an auxiliary 'hashtag prediction task', which was used as a projection into a Long Short Term Memory (LSTM) - Recurrent Neural Net (RNN) layer. *Wei et al. (2016)* and *Vijayaraghavan et al. (2016)* both implemented convolutional neural networks, however, did not achieve the same level of success as *Zarella et. al.*

There were also a handful of teams that took a 'featured-based' modelling approach. *Tutek et al. (2016)* used 'lexical features' such as n-grams, average word length, and number of hashtags as inputs to an ensemble classifier rooted inlogistic regression, gradient boosting machines, random forests, and support vector machines. Similarly, *(???)* used a voting classifier that took input from two multinomial Naive Bayes classifiers, one trained on word bi-grams while the other was trained on character tri-grams, and a logistic regression classifier trained on GloVe vectors. *Bohler et. al* also experimented with various additional features such as the presence or absence of negation and length of tweets.

Since the time the conference has ended, there have been two additional contributions to this task. Both *Sun et al. (2018)* and *Du et al. (2017)* leveraged neural attention networks to perform stance classification, with each team achieving higher final F1 scores than the winner of the initial competition. Although several of the competing teams found effective models, none were able to achieve higher final F1 scores than the SVM baseline *Mohammad et al. (2016b)*. This also holds for *Sun* and *Du*'s work after the conference ended. We will present our model's results in the context of performance against all SemEval-2016 Task 6 competing teams, *Mohammad et al. (2016b)*'s baseline, and the succeeding work by *Sun* and *Du*.

Outside the context of SemEval-2016 conference, the BERT language model from *Devlin et al. (2018)* has achieved state-of-the-art performance on several common natural language processing tasks including question answering, sentiment analysis, and named entity recognition [see *Figure 1* from *Devlin et al. (2018)*]. *Mayfield and Black (2019)* applied *Devlin et. al*'s BERT contextualized word embeddings to a stance classification task attempting to predict whether a Wikipedia user preferred to 'Keep' or 'Delete' a specific post based on their comments. *Mayfield et. al*'s results show that BERT out-performed their baseline classifiers, which used GloVe and Bag-of-Words embeddings as inputs.

Further, *Ma* experimented with various infrastructures built on top of BERT outputs for a classification task using Twitter data. The overall goal of *Ma*'s experiment was to determine whether a given tweet was 'on-topic' or 'off-topic' in the context of disaster management. *Ma* found that a bi-directional LSTM taking in the full BERT sequential output outperformed the standard BERT pooled output used for classification tasks.
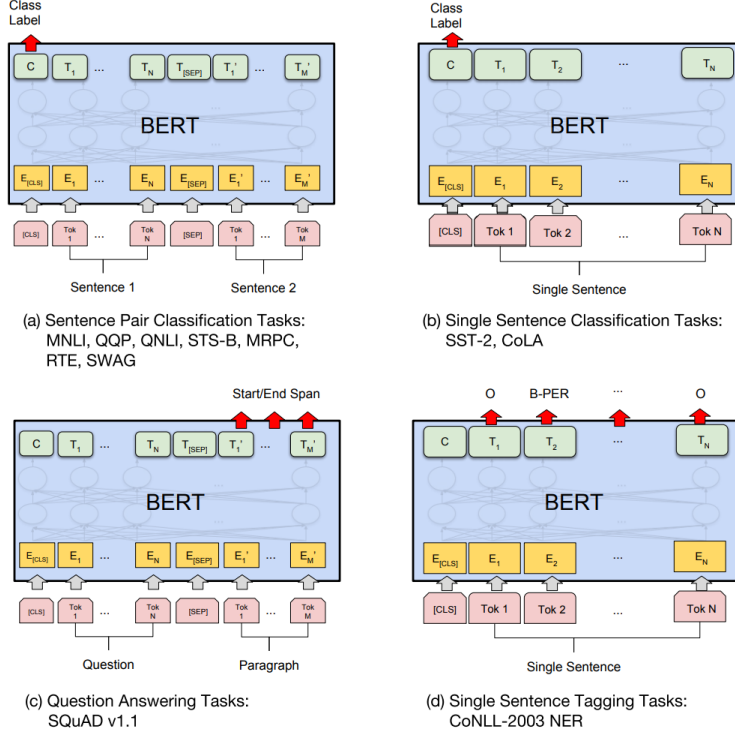
(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG

(b) Single Sentence Classification Tasks:
SST-2, CoLA

(c) Question Answering Tasks:
SQuAD v1.1

(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER

Figure 1: BERT by Task (Source: @devlin2018bert)

4 Model

Next, we describe the model we implemented to approach the SemEval-2016 Task 6 competition. After reviewing the work performed by contributing members of the competition, we hypothesized that a neural network which effectively captures the context of the limited characters and words contained in a tweet would perform well in classifying stance from Twitter data. Words contained in tweets cannot always be taken at face value. Hashtags, sarcasm, irony, and other forms of rhetoric are frequently found in tweets, and we would want our model to effectively capture these nuances since it contains valuable information about the author's stance.

4.1 Infrastructure

Given the state-of-the-art performance BERT has seen across many natural language processing tasks, including stance classification *[Mayfield]* and a classification task leveraging Twitter data *[Ma]*, we elected to leverage BERT for contextualized word embeddings. Rather than using the pooled output captured by BERT's '[CLS]' token, our model connects the full sequential output from BERT to a 128-unit LSTM layer. This infrastructure was motivated by the success of the BERT-LSTM integrated layers in *Ma et. al*'s Twitter classification task as well as *Zarella et. al*'s competition winning model that leveraged an auxiliary 'hashtag prediction task' as a projection into a LSTM-RNN layer. The LSTM layer output is fed into a densely connected layer consisting of 64 units using 'relu' activation function and He initialization. The densely connected layer feeds the classification layer which uses 'softmax' activation to predict the probability of each of the three stance labels in our data set: FAVOR/NONE/AGAINST. Dropout is performed at two points: 1) prior to the LSTM layer and 2) prior to the classification layer. Refer to *Figure 2* for a diagram of our model infrastructure.

As previously mentioned, we expect a model that has the capacity to capture the unique rhetoric seen in

Classification layer
(n x 3)

Densely connected
layer – 64 hidden
units
(n x 64)

LSTM layer – 128
hidden unit
(n x 128)

Bert contextualized
word embedding
(n x 83 x 768)

$T_{[CLS]}$ T_1 T_2 T_3 T_N

BERT

$E_{[CLS]}$ E_1 E_2 E_3 E_N

Tweet input tokens
(n x 83)
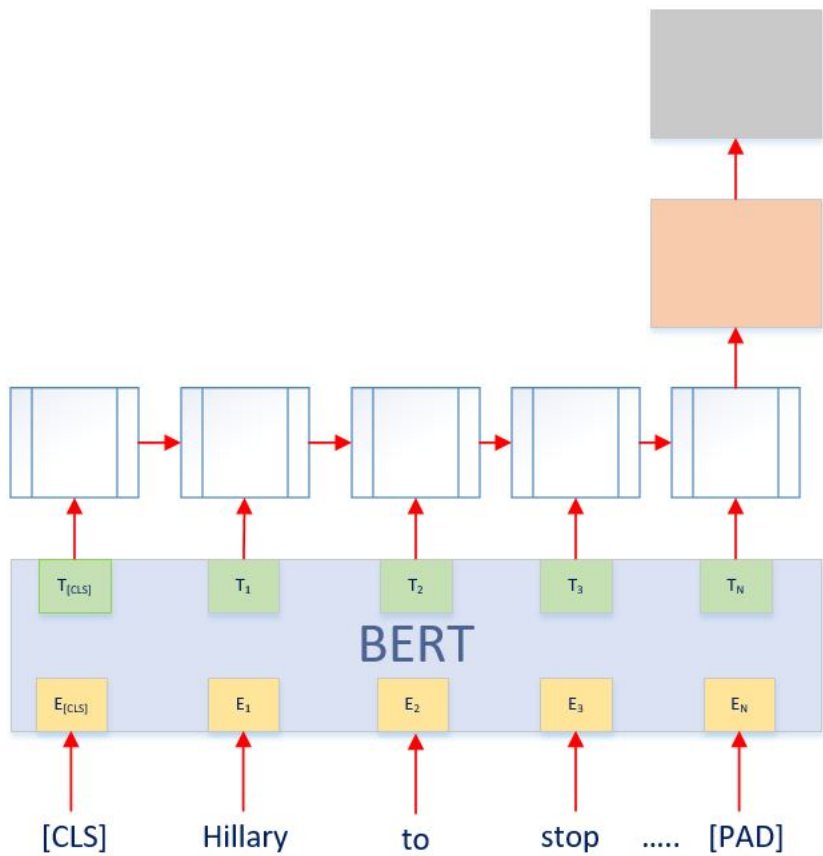
[CLS]   Hillary   to   stop   .....   [PAD]

Figure 2: Model Diagram

tweets to perform well in stance classification. We felt that the combined use of BERT with a LSTM layer would provide the infrastructure suitable to capture the nuances of hashtags, sarcasm, and irony frequently used on Twitter.

4.2 Hyperparameter Tuning

Since independent models were trained for each of the five topics, we tuned the model's hyperparameters on a single topic, Hillary Clinton, and then applied the results to the remaining topics. Hillary Clinton was selected because it contained the most consistent balance of classes within both the training and test sets. In addition, its majority class was the 'AGAINST' stance label, which was consistent with all other topics, excluding Climate Change. Our overall approach for hyperparameter tuning was trial and error. *Table 1* below details the hyperparameters we focused on as well as the final value implemented in our models.

Table 1: Hyperparameters tuned and final values used

Hyperparameter	Final Value
Learning Rate	0.001
Dropout Rate	50%
Batch Size	32
# of Bert Fine Tune Layers	6 layers
Tweet Pre-processing	None

Note we ran our model both with tweets pre-processed (all words lower-cased and punctuation and digits removed) as well as without pre-processing. The model performed consistently better without pre-processing applied, and thus, we elected to train all topic models on the raw tweets.

4.3 Training

As mentioned above, we leveraged a batch size of 32 tweets for training. Adam optimization with a learning rate of 0.001 was used with categorical cross entropy loss. We used 20 epochs for training on all topics, except for Climate Change. Climate Change had a considerably smaller sample size, and thus, required fewer epochs (7 epochs) before overfitting the model to the training data. Since no development data was provided, we carved out 15% of the training data to use as dev data to determine whether we were overfitting during training.

5 Results

Our model's results are summarized in *Table 2* and *Table 3* below *[Keep the ggplot tables? or redundant?]*. *Table 2* shows the final ranking of our model in the context of 1) the competing teams in the SemEval-2016 conference and 2) all contributions to the task, including the SVM baseline and post-conference work. *Table 3* shows our F1 scores for each topic in comparison to all teams that achieved a high score in at least 1 topic as well as the SVM baseline, *Sun et. al*, and *Du et. al*.

Table 2: Final F1 Scores with Ranks

Topic	F1 Score	SemEval-2016 Rank	Overall Rank
Overall F-macroT	60.88	1	2
Overall F-microT	69.72	1	3
Atheism	69.66	1	2

Topic	F1 Score	SemEval-2016 Rank	Overall Rank
Hillary Clinton	64.48	2	3
Abortion	64.06	1	3
Climate Change	48.47	3	5
Feminism	57.76	2	3

Table 3: F1 Scores for all state-of-the-art models (overall or in at least 1 target)

Team	Overall (F-microT)	Overall (F-macroT)	Atheism	Climate	Feminism	Hillary	Abortion
Muhammed et. al	70.32	59.01	69.19	43.80	58.72	61.74	66.91
Sun	69.79	61.00	70.53	49.56	57.50	61.84	66.16
Dessouky & Spittle	69.72	60.88	69.66	48.47	57.76	64.48	64.06
Du	68.79	59.56	59.77	53.59	55.77	65.38	63.72
MITRE	67.82	56.03	61.47	41.63	62.09	57.67	57.28
pkudblab	67.33	58.57	63.34	52.69	51.33	64.41	61.09
TakeLab	66.83	58.00	67.25	41.25	53.01	67.12	61.38
Majority Class	65.22	40.09	42.11	42.12	39.10	36.83	40.30
DeepStance	63.54	52.86	52.90	40.40	52.34	55.35	63.32
IDI@NTNU	62.47	55.08	59.59	54.86	48.59	57.89	54.47

The results in the above tables demonstrate the strength of our model. When comparing against teams participating in the SemEval2016 competition, we achieved the highest overall micro F1and macro F1 scores. We also achieved the top F1 score in the Atheism and Abortion topics, while finishing second in Hillary Clinton and Feminism, and third in Climate Change. When including *Sun et. al* and *Du et. al*'s post-conference contributions as well as the SVM baseline, we do not have any top scores. However, we still achieved the second highest macro F1 score, trailing *Sun et. al* by only 0.12. We also achieved the third highest micro F1 score, behind *Sun et. al* and *Muhammed et. al*. Note that the micro F1 score tends to be higher for models that perform better on the majority class, while macro F1 score tends to be higher for models that perform well on all classes. This can be seen in the comparison between the majority classifier's micro and macro F1 scores. The majority classifier scores a 65.22 on micro F1 score, but only 40.09 on macro F1. Our model achieved universally higher results than the majority classifier. See *Figure 4* below.

	Precision	Recall	F1.Score
Against	81.73	67.55	73.97
None	50.34	64.78	56.65
Favor	60.22	71.71	65.47

|	F1.Score
F−micro	69.72
F−macro	60.88
Atheism	69.66
Climate Change	48.47
Feminism	57.76
Hillary Clinton	64.48
Abortion	64.06

Figure 3: F-Score Tables

Figure 4: Results versus Majority Classifier

6 Limitations

While we were able to build an effective model for stance detection, we came across a number of apparent limitations. First, as mentioned within the **Data, Task, and Evaluation** section, the key query hashtags were removed from the training and test data, which would have been incredibly valuable in determining stance in the real world. Next, the stance labels were highly skewed. *Table 4* below shows the distribution of the labels for each target. While we allocated higher weights to the minority classes during the training process, our model still proved to be much more effective in predicting the majority class versus the minority classes.

Table 4: Data distribution

Target	# total	# train	favor (% of train)	none (% of train)	against (% of train)	# test	favor (% of test)	none (% of test)	against (% of test
Atheism	733	513	17.9	59.3	22.8	220	14.5	72.7	12.7
Climate Change	564	395	53.7	3.8	42.5	169	72.8	6.5	20.7
Feminist Movement	949	664	31.6	49.4	19.0	285	20.4	64.2	15.4
Hillary Clinton	984	689	17.1	57.0	25.8	295	15.3	58.3	26.4
Legalization of Abortion	933	653	18.5	54.4	27.1	280	16.4	57.3	18.4

In addition, there was a high level of disagreement amongst the label 'encoders', which can be translated to a higher expected error rate in our model. We can expect our model to struggle classifying the stance of tweets in which humans do not consistently agree fall under a specific stance label. Further, the total number of tweets in our dataset was low for a NLP problem. We felt that this may have affected the reproducibility of our model results. Re-running the same model on the identical training and test sets would produce slightly different results, despite fitting the training data to near 100% accuracy.

6.1 Error Analysis

See *Figure 1* for confusion matrix . . .

Errors for target **Atheism** were manually reviewed by the authors and categorized according to whether the purported "true" label or the predicted label from our model appeared to be correct (or if neither appeared correct).

It was surprising how many tweets appeared to have the incorrect true label associated with it. Some were blatantly and unambiguously wrong. Others has common challenges with the tweet text that may have caused some of these errors in annotation [*TO DO: add examples*]:

- *Referential*: Some tweets referenced very specific people or things that may be hard to align with the target or require a leap in assumption about connection of target (e.g. does supporting the NRA make you pro-choice, or are they just correlated)

- *Missing Context*:Some tweets were in reply to other tweets, whose content we cannot see, while others may suffer from missing the query hashtag

- *Sarcsm*: TBD

This apparent mis-labeling obviously causes some concern about the ability to train a model consistently that performs well on unseen data. However, there were additionally instances where our model made obvious
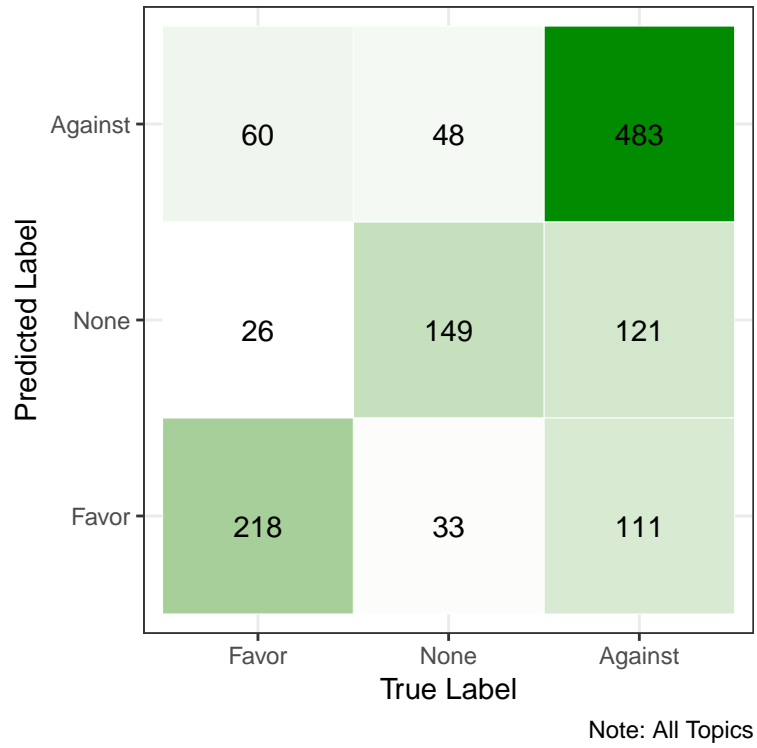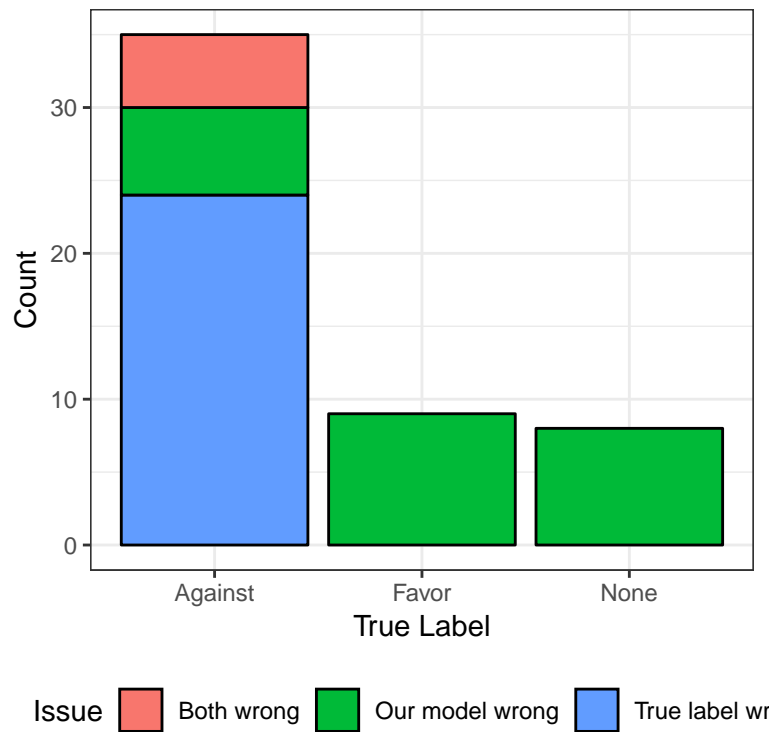
Figure 5: Confusion Matrix

Figure 6: Atheism Error Analysis Summary of Issues

mistakes. There were no systematic issues and it is difficult to determine what made these mis-classified tweets appear to have another stance. [**Not really sure what else to say here - reviewing errors not very enlightening**] It is notable that our model made consistent errors across true label categories, however the annotaters made errors exlusively in the against stance. The against stance constitutees 73% of test set data observations for the atheism target, indicating that some amount of the errors may be attributable to fatigue, perhaps, or some factor that causes repeated annotation to lean toward the majority class.

7 Conclusion and Future Work

In this paper, we implement BERT for contextualized word embeddings in a neural network model used to effectively detect stance. We propose a novel model architecture for use on the SemEval-2016 Task 6 dataset. While we did not beat the top score across all contributors to this task, we obtained results in line with the current state-of-the-art models. We achieved a F-microT of 69.72 and a F-macroT of 60.88, which ranked 3rd and 2nd, respectively, amongst all contributors to the task. This demonstrates the potential BERT has as a method of obtaining contextualized word embeddings for the purpose of stance classification.

Going forward, there are a couple additional areas in which we feel there is room for improvement. First, we would like to explore the use of a bidirectional LSTM layer in lieu of the current LSTM layer implemented in our model. In addition, we would like to investigate the benefits of using BERT Large in the embedding layer in place of BERT Base, which was utilized in our model.

Another potential area for improvement would be training BERT from scratch rather than leveraging the pre-trained weights and fine-tuning with the SemEval-2016 Task 6 data. If a large corpus could be obtained, similar to that described in the **Data, Task, and Evaluation** section above, training the Bert from scratch would likely result in a significant boost in F1 score.

Lastly, given the small data set size and skewness, putting effort into collecting more data to increase sample size and balance out the classes would likely benefit model performance.

References

Devlin, Jacob, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2018. "BERT: Pre-Training of Deep Bidirectional Transformers for Language Understanding."

Du, Jiachen, Ruifeng Xu, Yulan He, and Lin Gui. 2017. "Stance Classification with Target-Specific Neural Attention." In *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence, IJCAI-17*, 3988–94. doi:10.24963/ijcai.2017/557.

Mayfield, Elijah, and Alan Black. 2019. "Stance Classification, Outcome Prediction, and Impact Assessment: NLP Tasks for Studying Group Decision-Making." In *Proceedings of the Third Workshop on Natural Language Processing and Computational Social Science*, 65–77. Minneapolis, Minnesota: Association for Computational Linguistics. doi:10.18653/v1/W19-2108.

Mohammad, Saif, Svetlana Kiritchenko, Parinaz Sobhani, Xiaodan Zhu, and Colin Cherry. 2016a. "A Dataset for Detecting Stance in Tweets." In *Proceedings of the Tenth International Conference on Language Resources and Evaluation (LREC'16)*, 3945–52. Portorož, Slovenia: European Language Resources Association (ELRA). https://www.aclweb.org/anthology/L16-1623.

———. 2016b. "SemEval-2016 Task 6: Detecting Stance in Tweets." In *Proceedings of the 10th International Workshop on Semantic Evaluation (SemEval-2016)*, 31–41. San Diego, California: Association for Computational Linguistics. doi:10.18653/v1/S16-1003.

Sobhani, Parinaz, Saif Mohammad, and Svetlana Kiritchenko. 2016. "Detecting Stance in Tweets and Analyzing Its Interaction with Sentiment." In *Proceedings of the Fifth Joint Conference on Lexical and Computational*

Semantics, 159–69. Berlin, Germany: Association for Computational Linguistics. doi:10.18653/v1/S16-2021.

Sun, Qingying, Zhongqing Wang, Qiaoming Zhu, and Guodong Zhou. 2018. "Stance Detection with Hierarchical Attention Network." In *Proceedings of the 27th International Conference on Computational Linguistics*, 2399–2409. Santa Fe, New Mexico, USA: Association for Computational Linguistics. https://www.aclweb.org/anthology/C18-1203.

Tutek, Martin, Ivan Sekulić, Paula Gombar, Ivan Paljak, Filip Čulinović, Filip Boltužić, Mladen Karan, Domagoj Alagić, and Jan Šnajder. 2016. "TakeLab at SemEval-2016 Task 6: Stance Classification in Tweets Using a Genetic Algorithm Based Ensemble." In *Proceedings of the 10th International Workshop on Semantic Evaluation (SemEval-2016)*, 464–68. San Diego, California: Association for Computational Linguistics. doi:10.18653/v1/S16-1075.

Vijayaraghavan, Prashanth, Ivan Sysoev, Soroush Vosoughi, and Deb Roy. 2016. "DeepStance at SemEval-2016 Task 6: Detecting Stance in Tweets Using Character and Word-Level CNNs." In *Proceedings of the 10th International Workshop on Semantic Evaluation (SemEval-2016)*, 413–19. San Diego, California: Association for Computational Linguistics. doi:10.18653/v1/S16-1067.

Wei, Wan, Xiao Zhang, Xuqin Liu, Wei Chen, and Tengjiao Wang. 2016. "Pkudblab at SemEval-2016 Task 6 : A Specific Convolutional Neural Network System for Effective Stance Detection." In *Proceedings of the 10th International Workshop on Semantic Evaluation (SemEval-2016)*, 384–88. San Diego, California: Association for Computational Linguistics. doi:10.18653/v1/S16-1062.

Zarrella, Guido, and Amy Marsh. 2016. "MITRE at SemEval-2016 Task 6: Transfer Learning for Stance Detection." In *Proceedings of the 10th International Workshop on Semantic Evaluation (SemEval-2016)*, 458–63. San Diego, California: Association for Computational Linguistics. doi:10.18653/v1/S16-1074.